




influxdata[®]

The Modern Engine for Metrics and Events




Kapacitor

Introduction to Kapacitor

Agenda

- .. What is Kapacitor
- .. Understand TICKscript
 - .. Quoting rules
- .. Understand Batch vs Stream
- .. Use Kapacitor for
 - .. Downsampling
 - .. Alerting
- .. Creating Templated Tasks
- .. Using Topics
- .. Debugging TICKscript
- .. Examples
 - .. Working with Telegraf Data



Introduction to Kapacitor

What is it?

- .. Native time-series processing engine
- .. Process stream and batch data from InfluxDB
- .. Generates Alerts on dynamic criteria
 - .. match metrics for patterns
 - .. compute statistical anomalies
 - .. perform specific actions
- .. Integrates with 15+ Alert Providers
 - .. OpsGenie
 - .. Sensu
 - .. PagerDuty
 - .. Slack
 - .. and more.



Install



Installation

- .. Download
 - .. GitHub
 - .. <https://portal.influxdata.com/downloads>
- .. Pick your favorite package
- .. Install

Kapacitor Design

Server Daemon (kapacitord)

- Config for inputs/outputs/services/ etc.

Tasks - units of work

- Defined by a TICKscript
- Stream or Batch
- DAG - pipeline

Templates -

- Building many instances of a templated task

CLI (kapacitor)

- Calls for HTTP API or server
- Not interactive

Recordings - saved data

- Useful for isolated testing

Topics & topics handlers

- Generic rules for all alters



Defining Tasks



TICKScript

```
var measurement = 'requests'

var data = stream
  |from()
    .measurement(measurement)
  |where(lambda: "is_up" == TRUE)
  |where(lambda: "my_field" > 10)
  |window()
    .period(5m)
    .every(5m)

// Count number of points in window
data
  |count('value')
  .as('the_count')

// Compute mean of data window
data
  |mean('value')
  .as('the_average')
```

- Chain invocation language
 - | chains together different nodes
 - . refers to specific attributes on a node
- Variables refer to values
 - Strings
 - Ints, Floats
 - Durations
 - Pipelines

TICKScript Syntax - Quoting Rules

```
// ' means the use the literal string value
var measurement = 'requests'

var data = stream
  |from()
    .measurement(measurement)
// " means use the reference value
|where(lambda: "is_up" == TRUE)
|where(lambda: "my_field" > 10)
|window()
  .period(5m)
  .every(5m)

// ' means the use the literal string value
data
  |count('value')
    .as('the_count')
```

- Double Quotes
 - References data in lambda expression
- Single Quotes
 - Literal String value

Create a Stream TICKscript

```
// cpu.tick
stream
  |from()
    .measurement('cpu')
  |log()
```

- Logs all data from the measurement **cpu**

Create a Stream Task

```
// cpu.tick  
stream  
  |from()  
    .measurement('cpu')  
  |log()
```

```
$ kapacitor define cpu \  
-tick cpu.tick \  
-type stream \  
-dbrp telegraf.autogen
```

Create a Stream Task

```
// cpu.tick
dbrp "telegraf"."autogen"
dbrp "telegraf"."not_autogen"

stream
  |from()
    .measurement('cpu')
    .database('telegraf')
    .retentionPolicy('not_autogen')
  |log()
```

```
$ kapacitor define cpu \
  -tick cpu.tick \
```

Show a Stream Task

```
$ kapacitor define cpu \
  -tick cpu.tick \
  -type stream \
  -dbrp telegraf.autogen
```

```
$ kapacitor enable cpu
```

```
$ kapacitor show cpu
ID: cpu
Error:
Template:
Type: stream
Status: enabled
Executing: true
Created: 10 Oct 17 16:05 EDT
Modified: 10 Oct 17 16:05 EDT
LastEnabled: 01 Jan 01 00:00 UTC
Databases Retention Policies: ["telegraf"."autogen"]
TICKscript:
// cpu.tick
stream
  |from()
    .measurement('cpu')
  |log()

DOT:
digraph cpu {
stream0 -> from1;
from1 -> log2;
}
```

Create a More Interesting Stream TICKscript

```
// cpu.tick
stream
  |from()
    .measurement('cpu')
  |window()
    .period(5m)
    .every(1m)
  |mean('usage_user')
    .as('mean_usage_user')
  |log()
```

- Create 5m windows of data that emit every 1m
- Compute the average of the field **usage_user**
- Log the result

An even more interesting TICKscript

```
// cpu.tick
stream
  |from()
    .measurement('cpu')
  |where(lambda: "cpu" == 'cpu-total')
  |window()
    .period(5m)
    .every(1m)
  |mean('usage_user')
    .as('mean_usage_user')
  |log()
```

- Filter on the tag **cpu=cpu-total**
- Create 5m windows of data that emit every 1m
- Compute the average of the field **usage_user**
- Log the result

Adding Alerts

```
// cpu.tick
stream
  |from()
    .measurement('cpu')
  |where(lambda: "cpu" == 'cpu-total')
  |window()
    .period(5m)
    .every(1m)
  |mean('usage_user')
    .as('mean_usage_user')
  |alert()
    .crit(lambda: "mean_usage_user" > 80)
    .message('cpu usage high!')
    .slack()
      .channel('alerts')
    .email('oncall@example.com')
```

- Filter on the tag **cpu=cpu-total**
- Create 5m windows of data that emit every 1m
- Compute the average of the field **usage_user**
- Alert when **mean_usage_user > 80**
- Send alert to
 - Slack channel **alerts**
 - Email **oncall@example.com**

Create a Batch TICKscript

```
// batch_cpu.tick
batch
  |query(''
SELECT mean("usage_user") AS mean_usage_user
FROM "telegraf"."autogen"."cpu"
''')
    .period(5m)
    .every(1m)
  |log()
```

- Query 5m windows of data every 1m
- Compute the average of the field **usage_user**
- Log the result

Create a Batch Task

```
// batch_cpu.tick
batch
  |query(''
SELECT mean("usage_user") AS mean_usage_user
FROM "telegraf"."autogen"."cpu"
'')
    .period(5m)
    .every(1m)
  |log()
```

```
$ kapacitor define batch_cpu \
  -tick batch_cpu.tick \
  -type batch \
  -dbrp telegraf.autogen
```

Create a Batch TICKscript

```
// batch_cpu.tick
batch
  |query(''
SELECT mean("usage_user") AS mean_usage_user
FROM "telegraf"."autogen"."cpu"
''')
    .period(5m)
    .every(1m)
  |alert()
    .crit(lambda: "mean_usage_user" > 80)
    .message('cpu usage high!')
    .slack()
    .channel('alerts')
    .email('oncall@example.com')
```

- Query 5m windows of data every 1m
- Compute the average of the field **usage_user**
- Alert when **mean_usage_user > 80**
- Send alert to
 - Slack channel **alerts**
 - Email **oncall@example.com**

Batching

- Queries InfluxDB periodically
- Does not buffer as much of data in RAM
- Places additional query load on InfluxDB

Streaming

- Writes are mirrored onto the InfluxDB instance
- Buffers all data in RAM
- Places additional write load on Kapacitor




Using Kapacitor for Downsampling



What is Downsampling?


Downsampling is the process of reducing a sequence of points in a series to a single data point

- “ Computing the average, max, min, etc of a window of data for a particular series



Why Downsample?


- “ Faster queries
- “ Store less data



Downsample using Kapacitor

Offload computation to a separate host

- .. Create a task that aggregates your data
 - .. Can be stream or batch depending on your use-case
- .. Writes data back into InfluxDB
 - .. Typically back into another retention policy



Example

```
// batch_cpu.tick
batch
  |query(''
SELECT mean("usage_user") AS usage_user
FROM "telegraf"."autogen"."cpu"
''')
  .period(5m)
  .every(5m)
  |influxDBOut()
  .database('telegraf')
  .retentionPolicy('5m')
  .tag('source', 'kapacitor')
```

- Downsample the data into 5m windows
- Store that data back into a the **5m** retention policy in the **telegraf** database



Defining Task Templates

Create a Stream Template

```
// generic.tick
var measurement string
var where_filter = lambda: TRUE
var groups = [*]
var field string
var crit lambda
var window = 5m

stream
  |from()
    .measurement(measurement)
    .where(where_filter)
    .groupBy(groups)
  |window()
    .period(window)
    .every(window)
  |mean(field)
  |alert()
    .crit(crit)
    .slack()
    .channel(slack_channel)
```

```
$ kapacitor define-template generic \
  -tick generic.tick \
  -type stream \
```

Create a Stream Task from a Template

```
// generic_vars.json
```

```
{
  "measurement": {"type": "string", "value": "cpu" },
  "where_filter": {"type": "lambda", "value": "\"cpu\" == 'cpu-total'"},
  "groups": {"type": "list", "value": [{"type": "string", "value": "host"}, {"type": "string", "value": "dc"}]},
  "field": {"type": "string", "value": "usage_idle" },
  "crit": {"type": "lambda", "value": "\"mean\" < 10.0" },
  "window": {"type": "duration", "value": "1m" },
}
```

```
$ kapacitor define generic_task \
  -template generic \
  -vars generic_vars.json
  -dbrp telegraf.autogen
```

Show a Templated Task

```
stream
  |from()
    .measurement(measurement)
    .where(where_filter)
    .groupBy(groups)
  |window()
    .period(window)
    .every(window)
  |mean(field)
  |alert()
    .warn(warn)
    .crit(crit)
    .slack()
    .channel(slack_channel)
```

Vars:

Name	Type	Value
crit	lambda	"mean" < 10.0
field	string	usage_idle
groups	list	[host, dc]
measurement	string	cpu
warn	lambda	"mean" < 10.0
where_filter	lambda	"cpu" ==
'cpu-total'		
window	duration	1m0s


```
$ kapacitor show generic_task
```



TICKscript Examples


Example

- .. Compute the rolling average of the fields `usage_user`, `usage_system`, and `usage_idle`
- .. Create alerts for each




Example

- .. Compute the rolling average of the fields `usage_user`, `usage_system`, and `usage_system`
- .. Create alerts for each
- .. Join each of the computed values together




Example

- .. Create a derived metric under the measurement `cpu` that contains the `cpu usage` of a host
- .. Write the result back into a new database called `telegraf_kapacitor`




Example


- “ Change a field into a tag



Example

- Gracefully handle data that is missing fields





Advanced Kapacitor: Alert Topic Handlers

Michael DeSa

Agenda

- Describe what *topic* is as it relates to Kapacitor
- Use *topics* effectively in Ticksript
- Explain what a *topic handlers* are used for
- Define their own *topic handlers*
- Use special purpose *topic handlers*

Kapacitor Topics

- Specified in a Ticksript
- Allows users to separate the handling of alerts from the task that generates them
- Gives users the ability to handle alerts in a more sophisticated way
- Follows a publish/subscribe pattern
 - Alerts are published to a topic
 - Handlers subscribe to a topic
- Handlers are scoped to a topic

Without Topics

- All alert handlers had to be defined in tickscript
- Adding a new handler required redefining the task
 - All current task state would be lost
- Couldn't have cascading alerts be triggered
- Simplistic condition matching

```
// example.tick  
stream  
  |from()  
    .measurement('cpu')  
    .groupBy(*)  
  |alert()  
    .warn(lambda: "usage_idle" < 20)  
    .crit(lambda: "usage_idle" < 10)  
    .slack()  
      .channel('#alerts')  
    .slack()  
      .channel('#on_call')
```


With Topics: Creating/Using a topic

- Remove all handlers in the tickscript
- Add a topic handler

Note: Using a topic in a tickscript creates the topic.

```
// example.tick
stream
  |from()
    .measurement('cpu')
    .groupBy(*)
  |alert()
    .warn(lambda: "usage_idle" < 20)
    .crit(lambda: "usage_idle" < 10)
    .topic('basic')
```

With Topics: List topics

- Shows
 - Topic Name - ID
 - Current Alert Level - Level
 - Number of Alerts Published - Collected

```
$ kapacitor list topics
```

ID	Level	Collected
basic	OK	245

With Topics: Show Topic

- Shows
 - Topic Name - ID
 - Current Alert Level - Level
 - Number of Alerts Published - Collected
 - Sample of Collected Events - Events

```
$ kapacitor show-topic basic
ID: basic
Level: OK
Collected: 247
Handlers: []
Events:
Event Level Message Date
cpu OK cpu ... 05 Jul 17 20:55 EDT
cpu OK cpu ... 05 Jul 17 20:55 EDT
cpu OK cpu ... 05 Jul 17 20:55 EDT
cpu OK cpu ... 05 Jul 17 20:55 EDT
cpu OK cpu ... 05 Jul 17 20:55 EDT
cpu OK cpu ... 05 Jul 17 21:05 EDT
cpu OK cpu ... 05 Jul 17 21:05 EDT
cpu OK cpu ... 05 Jul 17 21:05 EDT
cpu OK cpu ... 05 Jul 17 21:10 EDT
```

Creating Topic Handlers

- Defining a topic handlers requires:
 - Topic Name
 - Handler Name
 - yaml File
- Yaml file specifies
 - Type of handler - kind
 - Typically the property method name from tickscript (e.g. slack, post, pagerduty, etc)
 - Required Parameters - options
 - Typically the property methods from tickscript

```
// alerts.yaml

kind: slack
topic: basic
id: alerts-channel
options:
  channel: "#alerts"
```

```
// on_call.yaml

kind: slack
topic: basic
id: on-call-chan
options:
  channel: "#on_call"
```

```
$ kapacitor define-topic-handler alerts.yaml
$ kapacitor define-topic-handler on_call.yaml
```

With Topics: List topics handlers

- Shows
 - Topic Name - Topic
 - Handler Name - ID
 - Type of handler - Kind

```
$ kapacitor list topic-handlers
```

Topic	ID	Kind
basic	alerts-channel	slack
basic	on-call-chan	slack

With Topics: Show Topic Handler

- Shows
 - Handler Name - ID
 - Handler Topic - Topic
 - Type of handler - Kind
 - Matching logic for alerts - Match
 - Handler Options - Options

```
$ kapacitor show-topic-handler basic  
alerts-channel  
ID: alerts-channel  
Topic: basic  
Kind: slack  
Match:  
Options: {"channel":"#alerts"}
```

```
$ kapacitor show-topic-handler basic  
on-call-chan  
ID: on-call-chan  
Topic: basic  
Kind: slack  
Match:  
Options: {"channel":"#on_call"}
```

With Topics: Show Topic (again)

- Shows
 - Topic Name - ID
 - Current Alert Level - Level
 - Number of Alerts Published - Collected
 - Current topic handlers - Handlers
 - Sample of Collected Events - Events

```
$ kapacitor show-topic basic
ID: basic
Level: OK
Collected: 285
Handlers: [alerts-channel, on-call-chan]
Events:
Event Level Message Date
cpu OK cpu ... 05 Jul 17 20:55 EDT
cpu OK cpu ... 05 Jul 17 20:55 EDT
cpu OK cpu ... 05 Jul 17 20:55 EDT
cpu OK cpu ... 05 Jul 17 20:55 EDT
cpu OK cpu ... 05 Jul 17 20:55 EDT
cpu OK cpu ... 05 Jul 17 21:05 EDT
cpu OK cpu ... 05 Jul 17 21:05 EDT
cpu OK cpu ... 05 Jul 17 21:05 EDT
cpu OK cpu ... 05 Jul 17 21:10 EDT
```

Chaining Topics using Handlers

- Topics can be chained together using a publish action.
 - This allows you to further group your alerts into various topics.

```
// chain.yaml
kind: publish
topic: basic
id: chain
options:
  topics:
    - ops_team
```

```
$ kapacitor define-topic-handler chain.yaml
```


Using Match Conditions in Handlers

- Conditions for matching a handle may be set in the match section of the yaml file
- Must be boolean expression
- Functions
 - Any built-in Kapacitor function
 - `changed()`
 - `level()`
 - `name()`
 - `taskName()`
 - `duration()`

```
// chain.yaml  
  
kind: publish  
topic: basic  
id: chain  
match: changed() == TRUE  
options:  
  topics:  
    - ops_team
```

```
$ kapacitor define-topic-handler chain.yaml
```

With Topics: Show Topic Handler

- Shows
 - Handler Name - ID
 - Handler Topic - Topic
 - Type of handler - Kind
 - Matching logic for alerts - Match
 - Handler Options - Options

```
$ kapacitor show-topic-handler basic chain
ID: chain
Topic: basic
Kind: publish
Match: changed() == TRUE
Options: {"topics":["ops_team"]}
```

Summary

- Kapacitor users can utilize **topics** instead of explicitly handling alerts in tickscript
 - Requires use of topic handlers
- Handlers are scoped to a topic
 - Typically referred to as **topic handlers**
- Topics may be chained together using the **publish** handler
- Handlers have additional matching logic to allow for more sophisticated alert event handling
- For more information visit the documentation
http://docs.influxdata.com/kapacitor/v1.3/guides/using_alert_topics/



Questions?





Thank you

